Julia

Feels like Python; Works like Lisp; Fast like C

Nick Eubank
CSDI, Vanderbilt University

April 16, 2018

Who am I?

- Post-Doc at Center for Study of Democratic Institutions
- Study social networks using cell-phone meta-data
  - Lots of simulations on networks with >10,000,000 nodes
- Regularly work with Stata, R, Python, and Julia
  - Some contributions to Julia packages, but I am *not* a core Julia developer!

Easy To Use Languages

Python, R, Matlab

Fast Languages

C, Java

Easy To Use Languages

Python, R, Matlab

- Interactive

Fast Languages

C, Java

- Compiled

Easy To Use Languages

Python, R, Matlab

- Interactive
- Dynamic typed

Fast Languages

C, Java

- Compiled
- Static Typed

Easy To Use Languages

Python, R, Matlab

- Interactive
- Dynamic typed
- Fast to write

Fast Languages

C, Java

- Compiled
- Static Typed
- Slow to write

Easy To Use Languages

Python, R, Matlab

- Interactive
- Dynamic typed
- Fast to write
- Slow to run

Fast Languages

C, Java

- Compiled
- Static Typed
- Slow to write
- Fast to run

- Write core libraries in C or Fortran
- Wrap in a nice language (R or Python)

Two Language Problem

Two Language Problem

- Hard to understand workings of packages

Two Language Problem

- Hard to understand workings of packages
- Hard to modify packages

Two Language Problem

- Hard to understand workings of packages
- Hard to modify packages
- Hard to write performant packages

Two Language Problem

- Hard to understand workings of packages
- Hard to modify packages
- Hard to write performant packages

$\Rightarrow$ True if you know C...

Two Language Problem

- Hard to understand workings of packages
- Hard to modify packages
- Hard to write performant packages

$\Rightarrow$ True if you know C...

$\Rightarrow$ *Extremely* true if you don't know C!

No Two Language Problem

Base Julia is written *in Julia*

- Even things like definitions of integers!

Most packages written in pure Julia

```python
# Python
def sum_sequence(start, stop):
    total = 0
    for i in range(start, stop):
        total = total + i
    return total

# Julia
function sum_sequence(start, stop)
    total = 0
    for i in start:stop
        total = total + i
    end
    return total
end
```

```python
# Python
def sum_sequence(start, stop):
    total = 0
    for i in range(start, stop):
        total = total + i
    return total
```

```julia
# Julia
function sum_sequence(start, stop)
    total = 0
    for i in start:stop
        total = total + i
    end
    return total
end
```

Python: `sum_sequence(0, 1000000)`: 78.8 milliseconds

Julia: `sum_sequence(0, 1_000_000)`: 0.0037 milliseconds

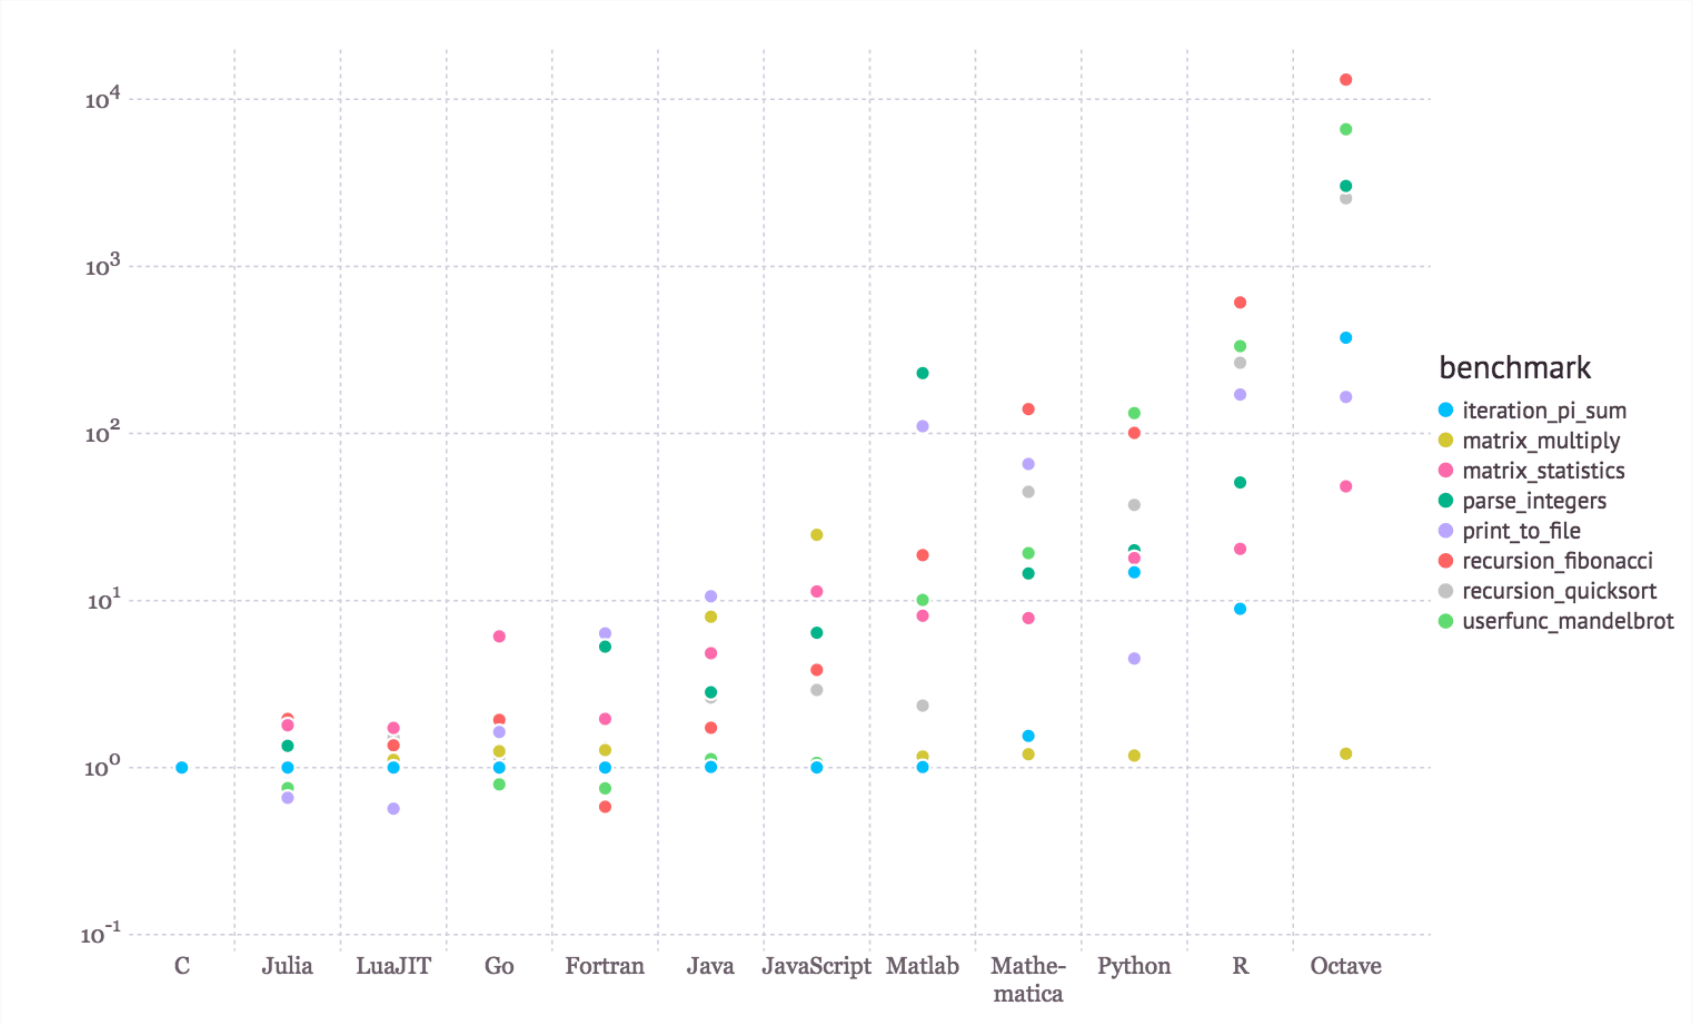
10^4
10^3
10^2
10^1
10^0
10^{-1}
C
Julia
LuaJIT
Go
Fortran
Java
JavaScript
Matlab
Mathe-
matica
Python
R
Octave
benchmark
iteration_pi_sum
matrix_multiply
matrix_statistics
parse_integers
print_to_file
recursion_fibonacci
recursion_quicksort
userfunc_mandelbrot

Why Python (and R) are slow

```python
# Python
def sum_sequence(0, 1000000):
    total = 0
    for i in range(start, stop):
        total = total + i
    return total
```

Why Python (and R) are slow

```python
# Python
def sum_sequence(0, 1000000):
    total = 0
    for i in range(start, stop):
        total = total + i
    return total
```

- Your processor doesn't know what "add total and i" means...

Why Python (and R) are slow

```
# Python
def sum_sequence(0, 1000000):
    total = 0
    for i in range(start, stop):
        total = total + i
    return total
```

- Your processor doesn't know what "add total and i" means...
  - Not all numbers are created equal

Why Python (and R) are slow

```
# Python
def sum_sequence(0, 1000000):
    total = 0
    for i in range(start, stop):
        total = total + i
    return total
```

- Your processor doesn't know what "add total and i" means...
  - Not all numbers are created equal
  - `+` actually has different meanings

Why Python (and R) are slow

```
# Python
def sum_sequence(0, 1000000):
    total = 0
    for i in range(start, stop):
        total = total + i
    return total
```

- Your processor doesn't know what "add total and i" means...
  - Not all numbers are created equal
  - `+` actually has different meanings

⇒Checks type of `total`, type of `i`, and looks up appropriate function `+` one million times!

Why Julia is Fast

```julia
# Julia
function sum_sequence(start, stop)
    total = 0
    for i in start:stop
        total = total + i
    end
    return total
end
```

Why Julia is Fast

```julia
# Julia
function sum_sequence(start, stop)
    total = 0
    for i in start:stop
        total = total + i
    end
    return total
end
```

- Treats function as a small program.

Why Julia is Fast

```
# Julia
function sum_sequence(start, stop)
    total = 0
    for i in start:stop
        total = total + i
    end
    return total
end
```

- Treats function as a small program.
- Realizes that `total` and `i` are always going to be integers, so only checks once.

Why Julia is Fast

```
# Julia
function sum_sequence(start, stop)
    total = 0
    for i in start:stop
        total = total + i
    end
    return total
end
```

- Treats function as a small program.
- Realizes that `total` and `i` are always going to be integers, so only checks once.
- Keeps copy of machine code once created so doesn't have to re-evaluate every time function is called.

Corollary: Julia is only fast inside functions

```
# Slow
total = 0
for i in 0:1_000_000
    total = total + i
end
```

Corollary: Julia is only fast inside functions

```
# Slow
total = 0
for i in 0:1_000_000
    total = total + i
end

# Fast
function sum_sequence(start, stop)
    total = 0
    for i in start:stop
        total = total + i
    end
    return total
end
sum_sequence(0, 1_000_000)
```

Features: Just Write the Loop

- No more need to always vectorize!
- But if you want, you still can with `.` notation.

```
x = rand(100)

# Loop
for i in 1:length(x)
    x[i] = sqrt(x[i])
end

# Vectorized
x = sqrt.(x)
```

Times: 6.651 ms (loop) and 7.682 ms (vectorized)

Features: Native Parallelism

Add workers:

```
addprocs(3)
```

Small jobs:

```
num_heads = @parallel (+) for i in 1:1_000_000
                rand(Bool)
            end
```

Or:

```
a = SharedArray{Float64}(1_000)
@parallel for i = 1:1_000
    a[i] = randn()
end
```

Features: Parallelism

Big jobs:

```
svds = pmap(svd, list_of_matrices)
```

Features: Support for Unicode

OLS with Unicode:

```
N = 4000
x = randn(N, 3)
ϵ = randn(N)
β = [2, 1, 90]
y = x * β + ϵ

β̂ = inv(x' * x) * x' * y
ϵ̂ = y - x * β̂
```

Features: Easy C Integration

If you need it, use `ccall`.

Features: Easy C Integration

If you need it, use `ccall`. Here's a call to `clock` function in C library `libc` that takes no arguments and returns an `Int32` value:

```
t = ccall((:clock, "libc"), Int32, ())
```

Features: Easy Python Integration

Import python `math` function and use its functions in Julia.

```
using PyCall
@pyimport math
math.sin(math.pi / 4) - sin(pi / 4)
```

Not 1.0 Yet...

Currently Stable Release: 0.6.2 Pending Release: 0.7

- Expected this summer ($\sim$ June 2018?)
- 0.7 is 1.0 with depreciation warnings
    - If you code works with 0.7, syntax won't change!

Expected changes

Expected changes

- Major compiler improvements for missing data

Expected changes

- Major compiler improvements for missing data
- New package manager

Expected changes

- Major compiler improvements for missing data
- New package manager
- Missing data type moving to core library

Hands-on Tutorials!

Go to `juliabox.com`, create an account, and navigate to `tutorials/intro-to-julia`.

Resources

Comparing Julia, Python, and Matlab side by side:

`https://cheatsheets.quantecon.org/`

Amazing cheatsheet for Julia:

`https://juliadocs.github.io/Julia-Cheat-Sheet/`