

Julia

Feels like Python; Works like Lisp; Fast like C

Nick Eubank

CSDI, Vanderbilt University

Resources at github.com/nickeubank/JuliaOverview

April 24, 2018

Show of Hands...

Work primarily in R?

Show of Hands...

Work primarily in Python?

Show of Hands...

Work primarily in Matlab?

Show of Hands...

Work regularly in C?

Show of Hands...

Social sciences?

Show of Hands...

Natural sciences?

Goals for Today

1. Why the need for *yet another* language?
2. Overview of Julia features
3. Brief hands-on tutorial
4. Leave you with resources for future exploration!

Who am I?

- Post-Doc at Center for Study of Democratic Institutions
- Study social networks using cell-phone meta-data
 - Lots of tabular data manipulations
 - Lots of simulations on networks with $>10,000,000$ nodes
- Regularly work with Stata, R, Python, and Julia
 - Some contributions to Julia packages, but I am *not* a core Julia developer!

Easy To Use Languages

Python, R, Matlab

Fast Languages

C, Java

Easy To Use Languages

Python, R, Matlab

- Interactive

Fast Languages

C, Java

- Compiled

Easy To Use Languages

Python, R, Matlab

- Interactive
- Dynamic typed

Fast Languages

C, Java

- Compiled
- Static Typed

Easy To Use Languages

Python, R, Matlab

- Interactive
- Dynamic typed
- Fast to write

Fast Languages

C, Java

- Compiled
- Static Typed
- Slow to write

Easy To Use Languages

Python, R, Matlab

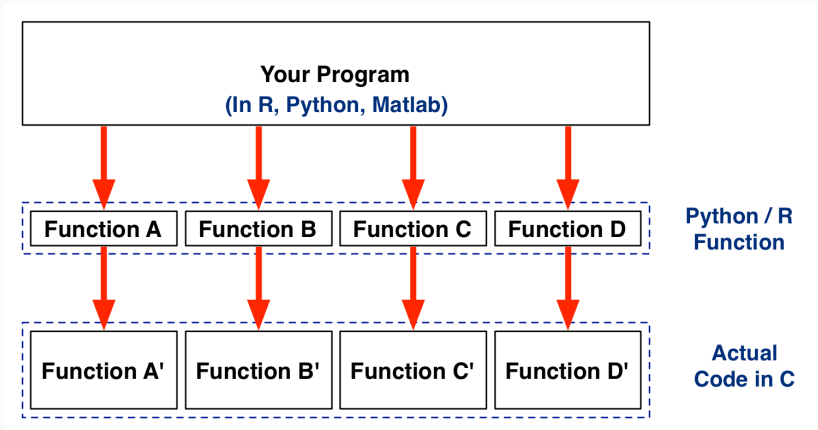
- Interactive
- Dynamic typed
- Fast to write
- Slow to run

Fast Languages

C, Java

- Compiled
- Static Typed
- Slow to write
- Fast to run

Hybrid Solution



Two Language Problem

Two Language Problem

- Hard to **understand** workings of packages

Two Language Problem

- Hard to **understand** workings of packages
- Hard to **modify** packages

Two Language Problem

- Hard to **understand** workings of packages
- Hard to **modify** packages
- Hard to **write** performant packages

Two Language Problem

- Hard to **understand** workings of packages
- Hard to **modify** packages
- Hard to **write** performant packages

⇒ True if you know C...

Two Language Problem

- Hard to **understand** workings of packages
- Hard to **modify** packages
- Hard to **write** performant packages

⇒ True if you know C...

⇒ *Extremely* true if you don't know C!

Julia: Solution Two Language Problem

Julia is a new, interactive, dynamic programming language written specifically with numerical computing in mind.

Base Julia is written *in Julia*

- Even things like definitions of integers!

Most packages written in pure Julia

Python

```
def sum_sequence(start, stop):  
    total = 0  
    for i in range(start, stop + 1):  
        total = total + i  
    return total
```

Julia

```
function sum_sequence(start, stop)  
    total = 0  
    for i in start:stop  
        total = total + i  
    end  
    return total  
end
```


Python

```
def sum_sequence(start, stop):  
    total = 0  
    for i in range(start, stop + 1):  
        total = total + i  
    return total
```


Julia

```
function sum_sequence(start, stop)  
    total = 0  
    for i in start:stop  
        total = total + i  
    end  
    return total  
end
```

Python: sum_sequence(0, 1000000): 78.8 milliseconds

R: sum_sequence(0, 1000000): 274 milliseconds

Julia: sum_sequence(0, 1_000_000): 0.0037 milliseconds



Why Python (and R) are slow

```
# Python
def sum_sequence(start, stop):
    total = 0
    for i in range(start, stop + 1):
        total = total + i
    return total
sum_sequence(0, 10000000)
```

Why Python (and R) are slow

```
# Python
```

```
def sum_sequence(start, stop):  
    total = 0  
    for i in range(start, stop + 1):  
        total = total + i  
    return total  
sum_sequence(0, 1000000)
```

- Your processor doesn't know what "add total and i" means...

Why Python (and R) are slow

```
# Python
```

```
def sum_sequence(start, stop):  
    total = 0  
    for i in range(start, stop + 1):  
        total = total + i  
    return total  
sum_sequence(0, 1000000)
```

- Your processor doesn't know what "add total and i" means...
 - Not all numbers are created equal

Why Python (and R) are slow

```
# Python
```

```
def sum_sequence(start, stop):  
    total = 0  
    for i in range(start, stop + 1):  
        total = total + i  
    return total  
sum_sequence(0, 1000000)
```

- Your processor doesn't know what "add total and i" means...
 - Not all numbers are created equal
 - + actually has different meanings

Why Python (and R) are slow

```
# Python
```

```
def sum_sequence(start, stop):  
    total = 0  
    for i in range(start, stop + 1):  
        total = total + i  
    return total  
sum_sequence(0, 1000000)
```

- Your processor doesn't know what "add total and i" means...
 - Not all numbers are created equal
 - + actually has different meanings

⇒ Checks type of **total**, type of **i**, and looks up appropriate function + one million times!

Why Julia is Fast

```
# Julia
function sum_sequence(start, stop)
    total = 0
    for i in start:stop
        total = total + i
    end
    return total
end
sum_sequence(0, 1_000_000)
```


Why Julia is Fast

```
# Julia
function sum_sequence(start, stop)
    total = 0
    for i in start:stop
        total = total + i
    end
    return total
end
sum_sequence(0, 1_000_000)
```

- Treats function as a small program.

Why Julia is Fast

```
# Julia
function sum_sequence(start, stop)
    total = 0
    for i in start:stop
        total = total + i
    end
    return total
end
sum_sequence(0, 1_000_000)
```

- Treats function as a small program.
- Realizes that `total` and `i` are always going to be integers, so only checks once.

Why Julia is Fast

```
# Julia
function sum_sequence(start, stop)
    total = 0
    for i in start:stop
        total = total + i
    end
    return total
end
sum_sequence(0, 1_000_000)
```

- Treats function as a small program.
- Realizes that `total` and `i` are always going to be integers, so only checks once.
- Keeps copy of machine code once created so doesn't have to re-evaluate every time function is called.

Corollary 1: Julia is only fast inside functions

```
# Slow  
total = 0  
for i in 0:1_000_000  
    total = total + i  
end
```

Corollary 1: Julia is only fast inside functions

Slow

```
total = 0
```

```
for i in 0:1_000_000
    total = total + i
end
```

Fast

```
function sum_sequence(start, stop)
    total = 0
    for i in start:stop
        total = total + i
    end
    return total
end
sum_sequence(0, 1_000_000)
```

Corollary 2: Type Stability for Max Speed

```
function return_if_even(a_number)
    if a_number % 2 == 0
        return a_number
    end
    if a_number % 2 != 0
        return "This is not even!"
    end
end
```

You can help the compiler by ensuring that, **conditional on the type of arguments**, all intermediate and output variables will always be of the same type.

Corollary 2: Type Stability for Max Speed

```
function return_if_even(a_number)
    if a_number % 2 == 0
        return a_number
    end
    if a_number % 2 != 0
        return "This is not even!"
    end
end
```

You can help the compiler by ensuring that, **conditional on the type of arguments**, all intermediate and output variables will always be of the same type.

This function is **not** type stable because:

- If `a_number` is an even integer, it returns an integer, but
- If `a_number` is an odd integer, it returns a string.

Features: Just Write the Loop

No more need to always vectorize!

```
x = rand(100)
```

```
# Loop  
for i in 1:length(x)  
    x[i] = sqrt(x[i])  
end
```

But you can if you want with `.` notation.

```
# Vectorized  
x = sqrt.(x)
```

Times: 6.651 ms (loop) and 7.682 ms (vectorized)

Features: Native Parallelism

Add workers:

```
addprocs(3)
```

Small jobs:

```
num_heads = @parallel (+) for i in 1:1_000_000  
    rand(Bool)  
end
```

Features: Native Parallelism

Add workers:

```
addprocs(3)
```

Small jobs:

```
num_heads = @parallel (+) for i in 1:1_000_000
    rand(Bool)
end
```

Or:

```
a = SharedArray{Float64}(1_000)
@parallel for i = 1:1_000
    a[i] = randn()
end
```

Features: Parallelism

Big jobs:

```
svds = pmap(svd, list_of_matrices)
```

Features: *Extensive* Linear Algebra Optimizations

```
julia> A = randn(n,n)
julia> Asym = A + A'
julia> issymmetric(Asym)
true
```

Features: *Extensive* Linear Algebra Optimizations

```
julia> A = randn(n,n)
julia> Asym = A + A'
julia> issymmetric(Asym)
true
```

Can also declare special structures (to deal with floating point errors):

- Triangular, Diagonal, Tridiagonal, Sparse Symmetric, etc..

Features: *Extensive* Linear Algebra Optimizations

```
julia> A = randn(n,n)
julia> Asym = A + A'
julia> issymmetric(Asym)
true
```

Can also declare special structures (to deal with floating point errors):

- Triangular, Diagonal, Tridiagonal, Sparse Symmetric, etc..

(Factorizations done using **LAPACK** and **UMFPACK** libraries)

Features: Scientific / Math Types Inbuilt

Base types:

- Rationals
- Imaginary Numbers
- BigInts

Features: Scientific / Math Types Inbuilt

Base types:

- Rationals
- Imaginary Numbers
- BigInts

Plus, user types as fast as Base types.

Features: Easy C Integration

If you need it, use `ccall`.

Features: Easy C Integration

If you need it, use `ccall`. Here's a call to `clock` function in C library `libc` that takes no arguments and returns an `Int32` value:

```
t = ccall{(:clock, "libc"), Int32, ()}
```

Features: Easy Python Integration

Import python `math` function and use its functions in Julia.

```
using PyCall
```

```
@pyimport math
```

```
math.sin(math.pi / 4) - sin(pi / 4)
```

Features: Support for Unicode

Ordinary least squares with Unicode:

```
N = 4000
```

```
x = randn(N, 3)
```

```
ε = randn(N)
```

```
β = [2, 1, 90]
```

```
y = x * β + ε
```

```
 $\hat{\beta} = \text{inv}(x' * x) * x' * y$ 
```

```
 $\hat{\epsilon} = y - x * \hat{\beta}$ 
```

Features: Meta-Programming

You can write Julia code that writes Julia code!

Contrasts with Python

Familiar:

- Duck-typing
- Pass by reference
- Iterators
- List (and array) comprehensions

Unfamiliar:

- No integer overflow checking
 - **SafeInts** package available
- Built in Package Manager
 - No name spaces yet; coming in new package manager.
- Not white-space sensitive
- Indexes start at 1, not 0
- Multiple dispatch for functions

Contrasts with R

Familiar:

- Multiple dispatch
- Built in package manager

Unfamiliar:

- No integer overflow checking
- Pass-by-reference and mutable / immutable data types
 - See [Eubank_PassByReference.pdf](#) on github.
- Loops as fast as vectorized functions

Not 1.0 Yet...

Currently Stable Release: 0.6.2

Pending Release: 0.7

- Expected this summer (~ June 2018?)
- 0.7 is 1.0 with depreciation warnings
 - If your code works with 0.7, syntax won't change!

Expected changes

- Handful of syntax changes
- Major compiler improvements for missing data
- New package manager

Hands-on Tutorials!

Go to `juliabox.com`, create an account, and navigate to `tutorials/intro-to-julia`.

Today we'll do:

- 6. Functions
- 10. Multiple Dispatch

Next Steps

For information on:

- How to install Julia
- How to get help with Julia
- Where to find more tutorials
- Most-used packages

See `Eubank_JuliaResources.pdf` at

www.github.com/nickeubank/JuliaOverview

